

```
In [1]: from matplotlib import pyplot as plt
import pandas as pd
import numpy as np
import torch
import time
import torch.optim as optim
from sklearn.preprocessing import MinMaxScaler, StandardScaler
```

```
In [2]: housing = pd.read_csv('Housing.csv')
housing
```

```
Out[2]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheat
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	
...
540	1820000	3000	2	1	1	yes	no	yes	
541	1767150	2400	3	1	1	no	no	no	
542	1750000	3620	2	1	1	yes	no	no	
543	1750000	2910	3	1	1	no	no	no	
544	1750000	3850	3	1	2	yes	no	no	

545 rows × 13 columns

```
In [3]: num_vars = ['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking']
df_new = housing[num_vars]
df_new
```

Out[3]:

	price	area	bedrooms	bathrooms	stories	parking
0	13300000	7420	4	2	3	2
1	12250000	8960	4	4	4	3
2	12250000	9960	3	2	2	2
3	12215000	7500	4	2	2	3
4	11410000	7420	4	1	2	2
...
540	1820000	3000	2	1	1	2
541	1767150	2400	3	1	1	0
542	1750000	3620	2	1	1	0
543	1750000	2910	3	1	1	0
544	1750000	3850	3	1	2	0

545 rows × 6 columns

In [4]: `df_new.shape`

Out[4]: (545, 6)

In [5]: `scaler = MinMaxScaler()
df_new[num_vars] = scaler.fit_transform(df_new[num_vars])
df_new.head()`

C:\Users\shenc\AppData\Local\Temp\ipykernel_46180\2943118898.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

`df_new[num_vars] = scaler.fit_transform(df_new[num_vars])`

Out[5]:

	price	area	bedrooms	bathrooms	stories	parking
0	1.000000	0.396564	0.6	0.333333	0.666667	0.666667
1	0.909091	0.502405	0.6	1.000000	1.000000	1.000000
2	0.909091	0.571134	0.4	0.333333	0.333333	0.666667
3	0.906061	0.402062	0.6	0.333333	0.333333	1.000000
4	0.836364	0.396564	0.6	0.000000	0.333333	0.666667

In [6]: `X = df_new.iloc[:, 1:6].values
Y = df_new.iloc[:, 0].values`

In [7]: `X = torch.tensor(X)
Y = torch.tensor(Y)`

```
In [8]: n_samples = X.shape[0]
        n_val = int(0.2 * n_samples)
        shuffled_indices = torch.randperm(n_samples)

        train_indices = shuffled_indices[:n_val]
        val_indices = shuffled_indices[n_val:]
```

```
In [9]: train_X = X[train_indices]
        train_Y = Y[train_indices]
        train_Y = train_Y.unsqueeze(1)

        val_X = X[val_indices]
        val_Y = Y[val_indices]
        val_Y = val_Y.unsqueeze(1)
```

```
In [10]: import torch
import torch.nn as nn
import torch.optim as optim

model = nn.Sequential(
    nn.Linear(5, 8),
    nn.ReLU(),
    nn.Linear(8, 1),
    nn.LogSoftmax(dim=1))
model = model.float()
learning_rate = 0.1

optimizer = optim.SGD(model.parameters(), lr=learning_rate)

out = model(train_X.float())

loss_fn = nn.MSELoss()

n_epochs = 200
correct = 0
total = 0

for epoch in range(n_epochs+1):
    training_start_time = time.time()
    out = model(train_X.float())
    loss = loss_fn(out, train_Y.float())

    out_v = model(val_X.float()) # <1>
    val_loss = loss_fn(out_v, val_Y.float())

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    if epoch==0 or epoch % 20 ==0:
        print("Epoch: %d, Training Loss: %f, Validation Loss: %f" % (epoch, float(loss),
_, predicted = torch.max(out, dim=1)
        total += val_Y.shape[0]
        correct += int((predicted == train_Y).sum())

print("Accuracy: %f" % (correct / total))
print('Time {:.2f}s'.format(time.time() - training_start_time))
```

```

Epoch: 0, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 20, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 40, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 60, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 80, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 100, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 120, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 140, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 160, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 180, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 200, Training Loss: 0.094215, Validation Loss: 0.095123
Accuracy: 8.000000
Time 0.00s

```

```

In [11]: import torch
import torch.nn as nn
import torch.optim as optim

model2 = nn.Sequential(
    nn.Linear(5,16),
    nn.ReLU(),
    nn.Linear(16,8),
    nn.ReLU(),
    nn.Linear(8,1),
    nn.ReLU(),
    nn.LogSoftmax(dim=1))
model = model.float()
learning_rate = 0.1

optimizer = optim.SGD(model.parameters(), lr=learning_rate)

out = model2(train_X.float())

loss_fn = nn.MSELoss()

n_epochs = 200
correct = 0
total = 0

for epoch in range(n_epochs+1):
    training_start_time = time.time()
    out = model2(train_X.float())
    loss = loss_fn(out, train_Y.float())

    out_v = model2(val_X.float()) # <1>
    val_loss = loss_fn(out_v, val_Y.float())

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    if epoch==0 or epoch % 20 ==0:
        print("Epoch: %d, Training Loss: %f, Validation Loss: %f" % (epoch, float(loss),
        _, predicted = torch.max(out, dim=1)
        total += val_Y.shape[0]
        correct += int((predicted == train_Y).sum()))

print("Accuracy: %f" % (correct / total))
print('Time {:.2f}s'.format(time.time() - training_start_time))

```

Epoch: 0, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 20, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 40, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 60, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 80, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 100, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 120, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 140, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 160, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 180, Training Loss: 0.094215, Validation Loss: 0.095123
Epoch: 200, Training Loss: 0.094215, Validation Loss: 0.095123
Accuracy: 8.000000
Time 0.00s